

Deployment Manual for Spring Boot Application using Kubernetes (k8s) and Docker

Contents

Prerequisites	1
Creating Package File of Spring Boot Application	1
Building Docker Image.....	1
Running Docker Container	2
Tagging and Pushing Docker Image	2
Deploying to Kubernetes	2
Verifying and Accessing the Application	3
Scaling and Replicating Pods in Kubernetes	4
Notes	4

Prerequisites

1. **Install Java Development Kit (JDK)**
 - Ensure you have JDK installed (openjdk:17-jdk-alpine). You can download it from [Oracle's JDK download page](#).
2. **Install Maven**
 - Maven is required to build the Spring Boot application. Download and install Maven from [Apache Maven's download page](#).
3. **Install Docker**
 - Docker is required to containerize the application. Install Docker from Docker's official website.
4. **Install Kubernetes and Minikube**
 - Minikube is used to run a local Kubernetes cluster. Install Minikube from Minikube's installation page.
5. **Install kubectl**
 - kubectl is a command-line tool for interacting with Kubernetes clusters. Install kubectl from Kubernetes' installation page.

Creating Package File of Spring Boot Application

1. **Clean and Package the Application**
 - Open your terminal or command prompt.
 - Navigate to the root directory of your Spring Boot application.
 - Run the following command to clean and package your application:

```
mvn clean package
```

- This will generate a JAR file in the `target` directory of your project.

Building Docker Image

2. **Build the Docker Image**
 - Ensure your Docker daemon is running.
 - In the root directory of your Spring Boot application, create a Dockerfile with the following content:

```
Dockerfile
FROM openjdk:17-jdk-alpine
EXPOSE 8080
ADD target/twit-app-0.0.1-SNAPSHOT.jar twit-app.jar
ENTRYPOINT ["java", "-jar", "/twit-app.jar"]
```

- Build the Docker image with the following command:

```
docker build -t twit-app:1.0 .
```

Running Docker Container

3. Run the Docker Container

- Start the Docker container and map port 8080 on your host to port 8080 in the container with the following command:

```
docker run -p 8080:8080 twit-app:1.0
```

- Your application should now be accessible at <http://localhost:8080>.

Tagging and Pushing Docker Image

4. Tag and Push the Docker Image

- Tag your Docker image for your Docker repository with the following command:

```
docker tag twit-app:1.0 <docker-id>/twit-app:1.0
```

- Push the Docker image to your repository:

```
docker push <docker-id>/twit-app:1.0
```

Deploying to Kubernetes

5. Apply Kubernetes Deployment

- Create a `deployment.yaml` file in the root directory with the following content:

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: twit-k8s
spec:
  replicas: 2
  selector:
    matchLabels:
      app: twit-app
  template:
    metadata:
      labels:
        app: twit-app
    spec:
      containers:
        - name: twit-app
          image: binodleo/twit-app:1.0
```

```
ports:
  - containerPort: 8080
```

- Apply the deployment configuration:

```
kubectl apply -f deployment.yaml
```

6. Apply Kubernetes Service

- Create a `service.yaml` file in the root directory with the following content:

```
apiVersion: v1
kind: Service
metadata:
  name: twit-k8s-service
spec:
  type: LoadBalancer # Exposes the service externally
                        # using a cloud provider's load balancer
  selector:
    app: twit-app # Matches the pods with this label
  ports:
    - protocol: TCP
      port: 80 # Port exposed by the service
      targetPort: 8080 # Port on the container to forward
                    # traffic to
```

- Apply the service configuration:

```
kubectl apply -f service.yaml
```

7. Load Image in Minikube

- Load the Docker image into Minikube:

```
minikube image load twit-app:1.0
```

Verifying and Accessing the Application

8. Verify Kubernetes Resources

- Check the status of all Kubernetes resources:

```
kubectl get all
```

- Ensure that the pods, services, and deployments are running as expected.

9. Access the Application

- Get the URL for your application service:

```
minikube service twit-app-service --url
```

- This command will provide the URL where your application is accessible.

10. Set Up Minikube Tunnel

- Start Minikube tunnel to expose services that use LoadBalancer type:

```
minikube tunnel
```

11. Open Minikube Dashboard

- Access the Minikube dashboard for a graphical view of your Kubernetes cluster:

```
minikube dashboard
```

Scaling and Replicating Pods in Kubernetes

12. Scaling Pods

- You can scale your application by increasing the number of replicas. To scale your deployment to 5 replicas, run

```
kubectl scale deployment twit-k8s --replicas=5
```

- Verify the scaling operation by checking the status of your pods:

```
kubectl get pods
```

13. Auto-scaling Pods

- Kubernetes supports horizontal pod auto-scaling based on resource utilization. To enable auto-scaling, use the following command:

```
kubectl autoscale deployment twit-k8s --min=2 --max=10 --cpu-percent=80
```

- This command configures the Horizontal Pod Autoscaler to maintain between 2 and 10 replicas of the application, scaling based on CPU utilization.

Notes

- **Ensure Docker Daemon is Running:** Always ensure that your Docker daemon is running before building and running Docker images.
- **Customize Configurations:** Customize the Dockerfile and Kubernetes YAML configurations as per your application's requirements.

- **Production Deployment:** For deploying to a production Kubernetes cluster, consider using a cloud provider like AWS EKS, Google Kubernetes Engine (GKE), or Azure Kubernetes Service (AKS).